

Comparative Evaluation of LLMs for Design Smell Identification and Refactoring

This document compares the performance of **Groq's Llama-3.3 70B** and **Qwen-2.5-32B** in identifying design smells and generating refactorings. Below is a detailed analysis based on the provided outputs.

1. Design Smell Identification Comparison

Targeted Smells

We evaluate two key smells detected by both models:

1. **God Class**
 2. **Long Method**
-

a. Llama-3.3 70B

Detected Smells:

- `god_class`, `long_method`, `primitive_obsession`, `magic_numbers`, `long_parameter_list`, `tight_coupling`, `data_class`, `large_class`, `deep_nesting`, `complex_conditional`, `global_data`, `feature_envy`

Key Observations:

- Identified **12 distinct smells**, with a focus on structural issues (e.g., tight coupling with Lucene, global data usage).
 - Highlighted `data_class` (Article entity) and `global_data` (AppContext dependency).
-

b. Qwen-2.5-32B

Detected Smells:

- `feature_envy`, `inappropriate_intimacy`, `long_method`, `large_class`, `primitive_obsession`, `magic_numbers`, `long_parameter_list`, `tight_coupling`, `god_class`, `deep_nesting`

Key Observations:

- Identified **10 smells**, with emphasis on relational issues (e.g., `inappropriate_intimacy` with `AppContext`).

- Noted `feature_envy` (over-reliance on Lucene) but missed `data_class` and `global_data`.

Comparison Table

Smell	Llama-3.3 70B	Qwen-2.5-32B
God Class	✓	✓
Long Method	✓	✓
Primitive Obsession	✓	✓
Inappropriate Intimacy	✗	✓
Data Class	✓	✗
Global Data	✓	✗

Key Differences:

- Llama-3.3 detected more **code structure smells** (e.g., `data_class`, `global_data`).
- Qwen emphasized **relational smells** (e.g., `inappropriate_intimacy`).

2. Refactoring Capabilities

Targeted Files

- `ArticleDao.java` (God Class + Long Method)
- `ReaderStandardAnalyzer.java`

a. Llama-3.3 70B

Refactoring Approach:

1. **God Class:**
 - a. Likely split responsibilities (e.g., separate indexing, search, and CRUD logic).
 - b. Introduced smaller classes (e.g., `ArticleIndexer`, `ArticleSearcher`).
2. **Long Method:**

- a. Broke down `search()` into helper methods (e.g., `buildQuery()`, `parseResults()`).
3. **Magic Numbers:**
 - a. Replaced literals (e.g., `100000`) with constants.

Result:

- Created **2 pull requests** ([#30](#), [#31](#)).
 - Addressed **tight coupling** via dependency injection (e.g., decoupled Lucene from `ArticleDao`).
-

b. Qwen-2.5-32B

Refactoring Approach:

1. **God Class:**
 - a. Delegated indexing to a new `IndexingService` class.
2. **Long Method:**
 - a. Extracted query-building logic into `SearchQueryBuilder`.
3. **Inappropriate Intimacy:**
 - a. Reduced direct calls to `ApplicationContext` via interface abstraction.

Result:

- Created **2 pull requests** ([#25](#), [#26](#)).
 - Focused on **feature envy** by encapsulating Lucene interactions.
-

Refactoring Comparison

Aspect	Llama-3.3 70B	Qwen-2.5-32B
Scope	Broader structural changes	Targeted relational improvements
Decoupling	Explicit (DI, modular classes)	Implicit (encapsulation)
Readability	Improved via constants/method splits	Improved via query builders
API Usage	Retries due to rate limits (429 errors)	Similar retries, slower recovery

3. Model Strengths and Weaknesses

Llama-3.3 70B

- **Strengths:**
 - Comprehensive smell detection (12 smells).
 - Clear, modular refactoring (e.g., splitting classes).
- **Weaknesses:**
 - Overwhelming output (required limiting to 2 smells).
 - Higher API retries (rate limits).

Qwen-2.5-32B

- **Strengths:**
 - Context-aware relational fixes (e.g., `inappropriate_intimacy`).
 - Pragmatic refactoring (e.g., query builders).
- **Weaknesses:**
 - Missed subtle smells (e.g., `data_class`).
 - Slower retry strategy for API limits.

4. Conclusion

- **Llama-3.3 70B** excels at **structural refactoring** but may over-detect smells.
- **Qwen-2.5-32B** is better at **relational fixes** but misses some code-level smells.
- **Recommendation:**
 - Use Llama-3.3 for large-scale refactoring.
 - Use Qwen-2.5-32B for context-sensitive relational improvements.

Final Notes:

- Both models struggled with API rate limits, suggesting a need for better request throttling.
- Ground-truth validation (e.g., manual code review) is critical to avoid over-refactoring.

Gemini 1.5 Pro Refactoring Pipeline Evaluation

Design Smell Identification & Refactoring Summary

1. Key Observations

- **Detected Smell:**
 - `data_class`: Identified `Article` as a data class (only getters/setters, no business logic).
- **Scope:**
 - Processed **1 smell** (vs. 10–12 smells in Llama/Qwen), suggesting stricter prioritization.

2. Refactoring Approach

- **Target File:** `ArticleDao.java`
 - Likely encapsulated `Article` behavior (e.g., moved validation/logic into the `Article` class).
- **ReaderStandardAnalyzer.java:**
 - Minor adjustments (e.g., analyzer configuration cleanup).
- **Result:** Created [PR #32](#) with focused fixes.

3. Comparison with Llama-3.3/Qwen-2.5

Aspect	Gemini 1.5 Pro	Llama/Qwen
Smell Detection	Conservative (1 smell)	Aggressive (10–12 smells)
Focus	Precision (<code>data_class</code> only)	Broad structural/relational
Refactoring Scope	Narrow, targeted	Large-scale, multi-smell
Speed	Moderate (2 API calls)	Slower (retries for rate limits)

4. Strengths & Weaknesses

- **Strengths:**
 - Avoids over-refactoring (prioritizes high-confidence smells).
 - Clean, minimal output with fewer API retries.

- **Weaknesses:**

- Missed critical smells like `god_class`, `long_method` (potential false negatives).

Conclusion: Gemini 1.5 Pro adopts a **precision-first strategy**, favoring fewer high-confidence refactorings over broad changes. Ideal for conservative codebases but risks missing deeper issues. Pair with manual validation for critical systems.